

# KPMG

July 27, 2020

```
[1]: import pandas as pd
import datetime as DT
import numpy as np
```

## 1 Transactions

Upon loading the first sheet of our dataset (Transactions) and reading its head, we see that the first row is just a notice. Hence, we skip it while reading. Then, we check its shape to see the number of rows (transaction entries) and columns (features).

Next, we check the number of null values in all columns. We observe that the column `online_order` has 360 null values. In order to ensure data completeness and quality, we can drop the rows with these null values. Depending on how we want to use our data later, we can also filter out rows on the basis of whether the order was online or not.

Similarly, we drop all rows that have null values for the brand name. After this, when we check null values again, we see that all rows with any null values have been dropped.

Finally, we change the `product_first_sold_date` from a vague string to a date format.

```
[2]: df = pd.read_excel("KPMG_VI_New_raw_data_update_final.xlsx", sheet_name=1,
    ↳ skiprows=1)

rows = df.shape[0]
cols = df.shape[1]
print("Number of Customer Id Entries: ", rows-1)
print("Number of Features: ", cols, "\n" )

# counting null values
print("Number of Initial Null Values: ")
print(df.isnull().sum(axis = 0), "\n")
# print(df.head(3))

# filtering out null values & offline orders
df = df[df['online_order'].notnull()]
df = df[df['brand'].notnull()]
# print(df.isnull().sum(axis = 0))
```

```
# converting string to date format
df['product_first_sold_date'] = pd.
    ↳to_datetime(df['product_first_sold_date'],unit='d',origin='1900-01-01')
df = df.reset_index(drop = True)

print(df.head(3))
```

Number of Customer Id Entries: 19999

Number of Features: 13

Number of Initial Null Values:

transaction_id	0
product_id	0
customer_id	0
transaction_date	0
online_order	360
order_status	0
brand	197
product_line	197
product_class	197
product_size	197
list_price	0
standard_cost	197
product_first_sold_date	197

dtype: int64

	transaction_id	product_id	customer_id	transaction_date	online_order	\
0	1	2	2950	2017-02-25	0.0	
1	2	3	3120	2017-05-21	1.0	
2	3	37	402	2017-10-16	0.0	

	order_status	brand	product_line	product_class	product_size	\
0	Approved	Solex	Standard	medium	medium	
1	Approved	Trek Bicycles	Standard	medium	large	
2	Approved	OHM Cycles	Standard	low	medium	

	list_price	standard_cost	product_first_sold_date
0	71.49	53.62	2012-12-04
1	2091.47	388.92	2014-03-05
2	1793.43	248.82	1999-07-22

## 1.1 Checking for Duplicates

In order to ensure uniqueness of data, we must ensure that there are no unnecessary duplicates. In this case, transaction ids should not be repeated over rows. Hence, we check for duplicate transaction ids and if any are found, we drop rows as required. Luckily, in this case, there are no duplicates.

```
[3]: print(len(df['transaction_id'])-len(df['transaction_id'].drop_duplicates()))
```

0

## 1.2 Net Profits

Since our ultimate goal is to help Sprocket Central Pty Ltd grow its business, we must consider their profits. For this, we can add a new column for net profits. This can help us achieve relevancy (data items with value meta data).

```
[4]: df['net_profit'] = df['list_price'] - df['standard_cost']
      print(df.net_profit.head(3))
```

0        17.87

1       1702.55

2       1544.61

Name: net\_profit, dtype: float64

Now that our dataset appears to be in a good shape, let's save our dataset into a dictionary that we will use later to generate an Excel file.

```
[5]: writer = pd.ExcelWriter('KPMG_TASK1.xlsx', engine='xlsxwriter')
      dataFrames = {'Transactions': df}
```

nhammad

## 2 Customer Demographics

Similar to the Transactions dataset, upon loading the third sheet (Customer Demographics), we see that the first row is just a notice. Hence, we skip it while reading it. The default columns seems quite absurd and useless. It's best to drop it in the early stage.

Then, we check its shape to see the number of rows and columns (features). Next, we check the number of null values in all columns. We observe that there are around 125 N/A values for the last name. However, since last names might not have any impact on our business strategy, we can choose to ignore that column. However, DOB is a valuable asset for us. Hence, we drop the rows which have null values for the DOB.

Similarly, job titles could also be useful for us. Hence, we drop the rows which have null values for the job title as well.

```
[6]: df = pd.read_excel("KPMG_VI_New_raw_data_update_final.xlsx", sheet_name=3,
      ↪ skiprows=1)
      df = df.drop(['default'], axis=1)

      rows = df.shape[0]
      cols = df.shape[1]
      print("Number of Customer Id Entries: ", rows-1)
      print("Number of Features: ", cols, "\n" )
```

```

print("Number of Initial Null Values: ")
print(df.isnull().sum(axis = 0), "\n")
df = df[df['DOB'].notnull()]
df = df[df['job_title'].notnull()]

print(df.head())

```

Number of Customer Id Entries: 3999

Number of Features: 12

Number of Initial Null Values:

```

customer_id          0
first_name           0
last_name           125
gender              0
past_3_years_bike_related_purchases  0
DOB                 87
job_title           506
job_industry_category  656
wealth_segment       0
deceased_indicator   0
owns_car             0
tenure              87
dtype: int64

```

	customer_id	first_name	last_name	gender	\
0	1	Laraine	Medendorp	F	
1	2	Eli	Bockman	Male	
2	3	Arlin	Dearle	Male	
4	5	Sheila-kathryn	Calton	Female	
7	8	Rod	Inder	Male	

	past_3_years_bike_related_purchases	DOB	job_title	\
0	93	1953-10-12	Executive Secretary	
1	81	1980-12-16	Administrative Officer	
2	61	1954-01-20	Recruiting Manager	
4	56	1977-05-13	Senior Editor	
7	31	1962-03-30	Media Manager I	

	job_industry_category	wealth_segment	deceased_indicator	owns_car	tenure
0	Health	Mass Customer	N	Yes	11.0
1	Financial Services	Mass Customer	N	Yes	16.0
2	Property	Mass Customer	N	Yes	15.0
4	NaN	Affluent Customer	N	Yes	8.0
7	NaN	Mass Customer	N	No	7.0

## 2.1 Contradiction in Genders

We notice that the values in the gender column are not consistent. By looking at the value counts, we observe that F and Female both represent the same thing. Similarly, M and Male represent the both thing. Additionally, terms U and Femal are also present in the dataset.

To fix this, we can replace Female/Femal with F and Male with M everywhere. Then, we can get rid of the row where gender is undefined (U).

```
[7]: print(df.gender.value_counts(), "\n")
df["gender"].replace({"Female": "F", "Male": "M", "Femal": "F"}, inplace=True)
df = df[df['gender'] != "U"]
```

```
Female    1769
Male      1643
U           1
F           1
Femal      1
M           1
Name: gender, dtype: int64
```

## 2.2 Customer Ages

When preparing business and marketing strategies, targeting the right age bracket is quite essential. Hence, we need to know the exact ages of our customers. For this, we can add a new column for customer ages. We can calculate ages by subtracting the DOB from the current date. After this, the dataset seems to look good to go.

```
[8]: now = pd.Timestamp('now')
df['age'] = (now - df['DOB']).astype('<m8[Y]')
df['age'] = df['age'].astype(np.int64)
print(df.head(3))
```

```
customer_id first_name last_name gender \
0           1   Laraine  Medendorp    F
1           2     Eli    Bockman     M
2           3    Arlin   Dearle      M

past_3_years_bike_related_purchases  DOB  job_title \
0                                93 1953-10-12  Executive Secretary
1                                81 1980-12-16  Administrative Officer
2                                61 1954-01-20    Recruiting Manager

job_industry_category wealth_segment deceased_indicator owns_car  tenure \
0           Health  Mass Customer                N      Yes    11.0
1  Financial Services  Mass Customer                N      Yes    16.0
2           Property  Mass Customer                N      Yes    15.0
```

```

    age
0    66
1    39
2    66

```

```
[9]: dataFrames['CustomerDemographic'] = df
```

### 3 Customer Addresses

Just like the previous two datasets, this one also has a notice in the first row. Hence, we skip it while loading our dataset. Then, we check its shape of our dataset to see the number of rows and columns (features).

Next, we check the number of null values in all columns. Luckily, there are none.

```
[10]: df = pd.read_excel("KPMG_VI_New_raw_data_update_final.xlsx", sheet_name=4,
    ↪ skiprows=1)

rows = df.shape[0]
cols = df.shape[1]
print("Number of Customer Id Entries: ", rows-1)
print("Number of Features: ", cols, "\n" )

print("Number of Initial Null Values: ")
print(df.isnull().sum(axis = 0), "\n")
print(df.head())
```

Number of Customer Id Entries: 3998

Number of Features: 6

Number of Initial Null Values:

```

customer_id      0
address          0
postcode         0
state            0
country          0
property_valuation 0
dtype: int64

```

	customer_id	address	postcode	state	country \
0	1	060 Morning Avenue	2016	New South Wales	Australia
1	2	6 Meadow Vale Court	2153	New South Wales	Australia
2	4	0 Holy Cross Court	4211	QLD	Australia
3	5	17979 Del Mar Point	2448	New South Wales	Australia
4	6	9 Oakridge Court	3216	VIC	Australia

	property_valuation
0	10
1	10
2	9
3	4
4	9

### 3.1 Contradiction in States

Upon seeing the head of our dataset, we observe something strange. Some state names are written full while others are abbreviations. To investigate further, we can check the value counts. After checking the value counts, we realize that there's a contradiction in state names. For instance, New South Wales and NSW is the same but are written separately. Same is the case with Victoria and VIC. Hence, we make the necessary replacements.

```
[11]: print(df.state.value_counts(), "\n")
df["state"].replace({"New South Wales": "NSW", "Victoria": "VIC"}, inplace=True)
print(df.head())
```

NSW	2054
VIC	939
QLD	838
New South Wales	86
Victoria	82

Name: state, dtype: int64

	customer_id	address	postcode	state	country \
0	1	060 Morning Avenue	2016	NSW	Australia
1	2	6 Meadow Vale Court	2153	NSW	Australia
2	4	0 Holy Cross Court	4211	QLD	Australia
3	5	17979 Del Mar Point	2448	NSW	Australia
4	6	9 Oakridge Court	3216	VIC	Australia

	property_valuation
0	10
1	10
2	9
3	4
4	9

Everything appears to be better now. As the final step, let's merge the datasets into a single Excel sheet.

```
[12]: dataFrames['CustomerAddress'] = df
for sheet, frame in dataFrames.items():
    frame.to_excel(writer, sheet_name = sheet, index=False)
writer.save()
```