

KPMG TASK 2

August 6, 2020

```
[1]: import pandas as pd
import datetime as DT
import numpy as np
import matplotlib.pyplot as plt
```

1 Customer Demographics

Upon loading the dataset, we see that the first row is just a notice. Hence, we skip it while reading it.

Then, we check its shape to see the number of rows and columns (features). Next, we check if there are duplicates for the customer id field. Luckily, there are none.

The default columns seems quite absurd and useless. It's best to drop it in the early stage.

```
[2]: df = pd.read_excel("KPMG_VI_New_raw_data_update_final.xlsx", sheet_name=3,
    ↳ skiprows=1)

rows = df.shape[0]
cols = df.shape[1]
print("Number of Customer Id Entries: ", rows)
print("Number of Features: ", cols, "\n" )

print("Number of Duplicates", len(df['customer_id'])-len(df['customer_id'].
    ↳ drop_duplicates()))

df = df.drop(['default'], axis=1)
df.head()
```

Number of Customer Id Entries: 4000

Number of Features: 13

Number of Duplicates 0

```
[2]:   customer_id  first_name last_name gender \
0           1      Laraine  Medendorp      F
1           2         Eli    Bockman    Male
```

2	3	Arlin	Dearle	Male
3	4	Talbot	NaN	Male
4	5	Sheila-kathryn	Calton	Female

	past_3_years_bike_related_purchases	DOB	job_title \
0	93	1953-10-12	Executive Secretary
1	81	1980-12-16	Administrative Officer
2	61	1954-01-20	Recruiting Manager
3	33	1961-10-03	NaN
4	56	1977-05-13	Senior Editor

	job_industry_category	wealth_segment	deceased_indicator	owns_car	tenure
0	Health	Mass Customer		N	Yes 11.0
1	Financial Services	Mass Customer		N	Yes 16.0
2	Property	Mass Customer		N	Yes 15.0
3	IT	Mass Customer		N	No 7.0
4	NaN	Affluent Customer		N	Yes 8.0

1.1 Checking for Null Values

Now, we check the number of null values in all columns. We observe that there are around 125 N/A values for the last name. However, since last names might not have any impact on our business strategy, we can choose to ignore that column. However, DOB is a valuable asset for us. Hence, we drop the rows which have null values for the DOB.

Similarly, job titles could also be useful for us. Hence, we drop the rows which have null values for the job title as well.

```
[3]: print("Number of Initial Null Values: ")
      print(df.isnull().sum(axis = 0),"\n")

      df = df[df['DOB'].notnull()]
      df = df[df['job_title'].notnull()]
```

```
Number of Initial Null Values:
customer_id           0
first_name            0
last_name            125
gender               0
past_3_years_bike_related_purchases  0
DOB                  87
job_title            506
job_industry_category  656
wealth_segment        0
deceased_indicator    0
owns_car              0
tenure               87
```

```
dtype: int64
```

1.2 Customer Ages

When preparing business and marketing strategies, targeting the right age bracket is quite essential. Hence, we need to know the exact ages of our customers. For this, we can add a new column for customer ages. We can calculate ages by subtracting the DOB from the current date.

```
[4]: now = pd.Timestamp('now')
df['age'] = (now - df['DOB']).astype('<m8[Y]')
df['age'] = df['age'].astype(np.int64)
df = df.reset_index(drop = True)
df.head()
```

```
[4]:
```

	customer_id	first_name	last_name	gender	\
0	1	Laraine	Medendorp	F	
1	2	Eli	Bockman	Male	
2	3	Arlin	Dearle	Male	
3	5	Sheila-kathryn	Calton	Female	
4	8	Rod	Inder	Male	

	past_3_years_bike_related_purchases	DOB	job_title	\
0	93	1953-10-12	Executive Secretary	
1	81	1980-12-16	Administrative Officer	
2	61	1954-01-20	Recruiting Manager	
3	56	1977-05-13	Senior Editor	
4	31	1962-03-30	Media Manager I	

	job_industry_category	wealth_segment	deceased_indicator	owns_car	\
0	Health	Mass Customer	N	Yes	
1	Financial Services	Mass Customer	N	Yes	
2	Property	Mass Customer	N	Yes	
3	NaN	Affluent Customer	N	Yes	
4	NaN	Mass Customer	N	No	

	tenure	age
0	11.0	66
1	16.0	39
2	15.0	66
3	8.0	43
4	7.0	58

1.3 Removing Contradiction in Genders

We notice that the values in the gender column are not consistent. By looking at the value counts, we observe that F and Female both represent the same thing. Similarly, M and Male represent the same thing. Additionally, terms U and Femal are also present in the dataset.

To fix this, we can replace Female/Femal with F and Male with M everywhere. Then, we can get rid of the row where gender is undefined (U). After this, the dataset seems to look good to go.

```
[5]: print(df.gender.value_counts(), "\n")
df["gender"].replace({"Female": "F", "Male": "M", "Femal": "F"}, inplace=True)
```

```
Female    1769
Male      1643
F           1
M           1
U           1
Femal       1
Name: gender, dtype: int64
```

1.4 Visualizing Purchases Based on Gender

First we plot a bar chart of number of purchases according to different genders. Since the value for U is really small, we can drop that value and then visualize again.

```
[6]: total = df.gender.size
females = (df.gender == "F").sum()
males = (df.gender == "M").sum()
unknown = (df.gender == "U").sum()

gender = ['Male', 'Female', 'Unknown']
numbers = [males, females, unknown]

plt.figure(figsize=(8,8))
bars = plt.bar(gender, numbers, width=0.4, bottom=None, align='center',
               data=None, color='cornflowerblue')
plt.xlim(-0.9, len(gender) - 1 + 0.9)
for i in range(len(numbers)):
    percentage = ((numbers[i]/total)*100)
    percentage = str(round(percentage,2)) + " %"
    plt.annotate(percentage, xy=(gender[i], numbers[i] + 10), ha='center')
plt.xlabel('Gender')
plt.ylabel('Number of Customers')
plt.title('Gender Distribution of Customers')
plt.show()
```

```

df = df[df['gender'] != "U"]

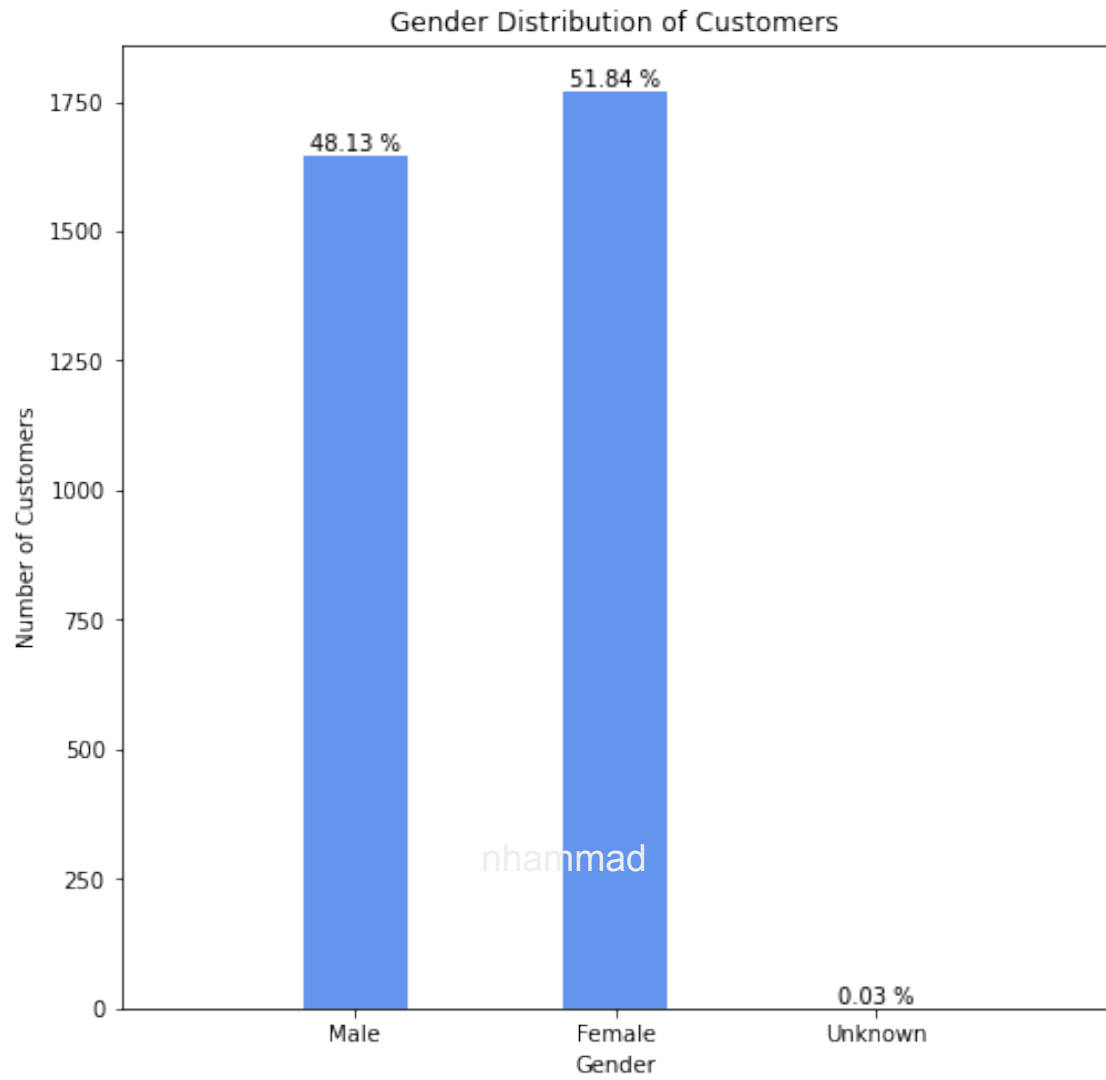
females = (df.gender == "F").sum()
males = (df.gender == "M").sum()

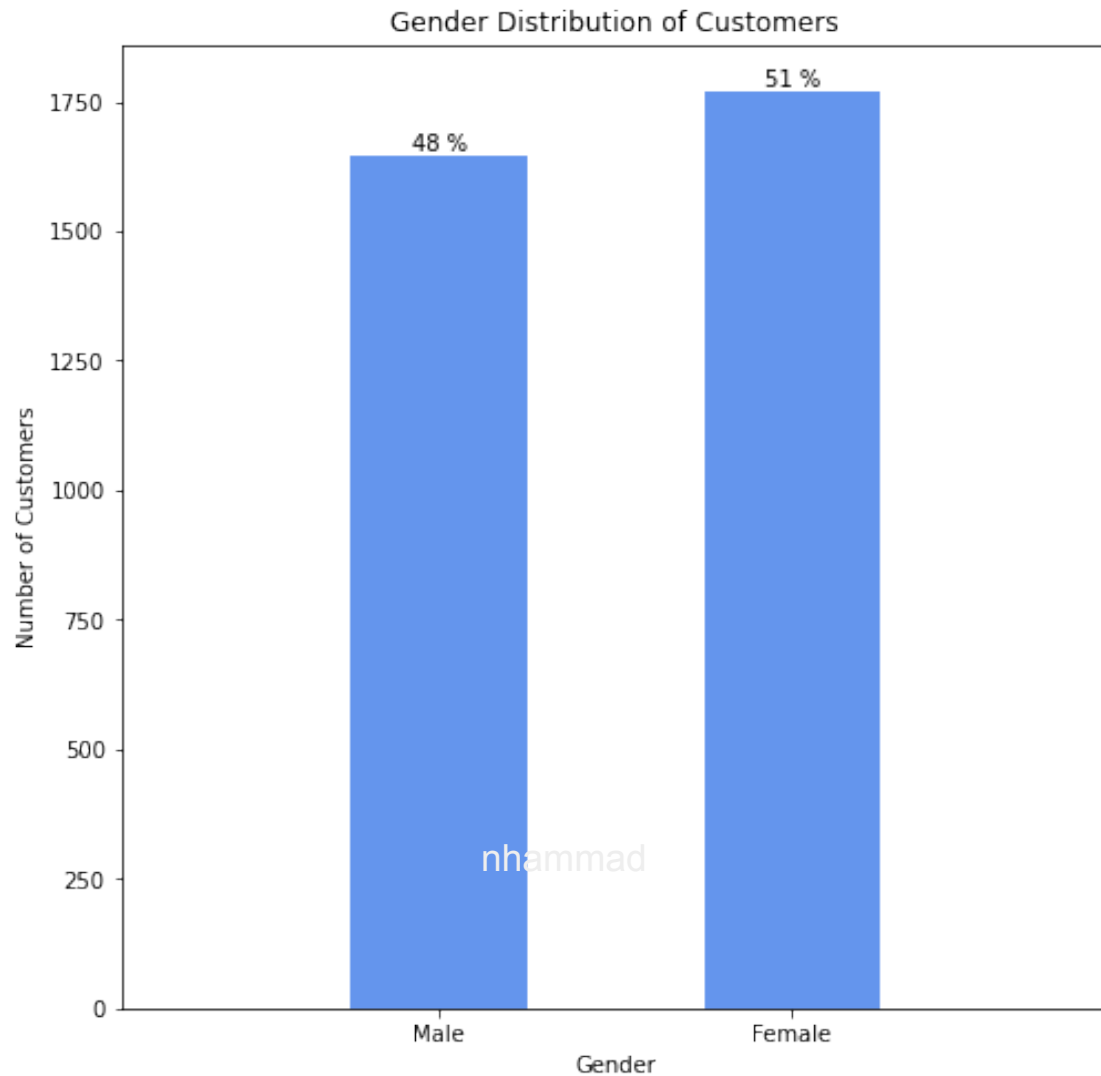
gender = ['Male', 'Female']
numbers = [males, females]

plt.figure(figsize=(8,8))
bars = plt.bar(gender, numbers, width=0.5, bottom=None, align='center',
    ↳data=None, color='cornflowerblue')
plt.xlim(-0.9, len(gender) - 1 + 0.9)
for i in range(len(numbers)):
    percentage = ((numbers[i]/(total-1))*100)
    percentage = str(int(percentage)) + " %"
    plt.annotate(percentage, xy=(gender[i], numbers[i] + 10), ha='center')
plt.xlabel('Gender')
plt.ylabel('Number of Customers')
plt.title('Gender Distribution of Customers')
plt.show()

```

nhammad





1.5 Visualizing Purchases Based on Age Groups

First we check the average age. Then, we plot a bar chart to see how the number of purchases are distributed amongst different age groups.

```
[7]: print("Average Age" , int(df.age.mean()))
```

```
under20 = 0  
_20To30 = 0  
_30To40 = 0  
_40To50 = 0  
_50To60 = 0  
_60To70 = 0
```

```

_70To80 = 0
above80 = 0

size = df.gender.size

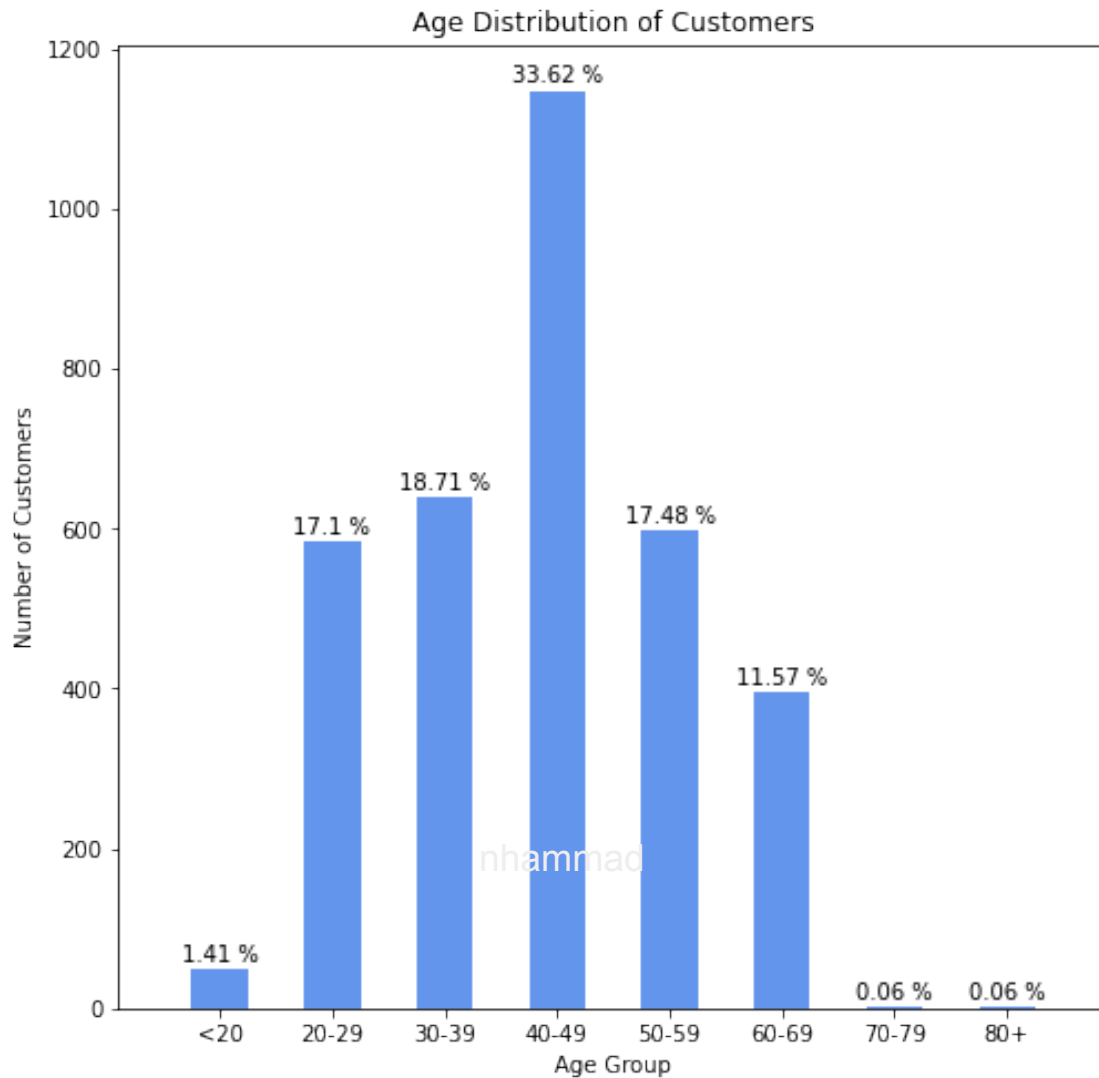
for value in df['age']:
    if(value<20):
        under20+=1
    elif (value >=20 and value <30):
        _20To30 += 1
    elif (value >= 30 and value <40):
        _30To40 += 1
    elif (value >= 40 and value <50):
        _40To50 += 1
    elif (value >= 50 and value <60):
        _50To60 += 1
    elif (value >= 60 and value <70):
        _60To70 += 1
    elif (value >= 70 and value <80):
        _70To80 += 1
    else:
        above80 += 1

ageGroups = ['<20', '20-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80+']
numbers = [under20, _20To30, _30To40, _40To50, _50To60, _60To70, _70To80,
↪above80]

plt.figure(figsize=(8,8))
bars = plt.bar(ageGroups, numbers, width=0.5, bottom=None, align='center',
↪data=None, color='cornflowerblue')
plt.xlim(-0.9, len(ageGroups) - 1 + 0.9)
for i in range(len(numbers)):
    percentage = ((numbers[i]/(size))*100)
    percentage = str(round(percentage,2)) + " %"
    plt.annotate(percentage, xy=(ageGroups[i], numbers[i] + 10), ha='center')
plt.xlabel('Age Group')
plt.ylabel('Number of Customers')
plt.title('Age Distribution of Customers')
plt.show()

```

Average Age 42



1.6 Visualizing Purchases Based on Gender & Age Groups

```
[8]: femalesUnder20 = 0
females20To30 = 0
females30To40 = 0
females40To50 = 0
females50To60 = 0
females60To70 = 0
females70To80 = 0
femalesAbove80 = 0

femaleAgesData = df[df['gender']=="F"]
```

```

for value in femaleAgesData['age']:
    if(value<20):
        femalesUnder20+=1
    elif (value >=20 and value <30):
        females20To30 += 1
    elif (value >= 30 and value <40):
        females30To40 += 1
    elif (value >= 40 and value <50):
        females40To50 += 1
    elif (value >= 50 and value <60):
        females50To60 += 1
    elif (value >= 60 and value <70):
        females60To70 += 1
    elif (value >= 70 and value <80):
        females70To80 += 1
    else:
        femalesAbove80 += 1

```

```

malesUnder20 = 0
males20To30 = 0
males30To40 = 0
males40To50 = 0
males50To60 = 0
males60To70 = 0
males70To80 = 0
malesAbove80 = 0

```

nhammad

```

maleAgesData = df[df['gender']=="M"]

```

```

for value in maleAgesData['age']:
    if(value<20):
        malesUnder20+=1
    elif (value >=20 and value <30):
        males20To30 += 1
    elif (value >= 30 and value <40):
        males30To40 += 1
    elif (value >= 40 and value <50):
        males40To50 += 1
    elif (value >= 50 and value <60):
        males50To60 += 1
    elif (value >= 60 and value <70):
        males60To70 += 1
    elif (value >= 70 and value <80):
        males70To80 += 1
    else:

```

```

    malesAbove80 += 1

N = 8
labels = ['<20', '20-29', '30-39', '40-49', '50-59', '60-69', '70-79', '80+']
maleAges = (malesUnder20, males20To30, males30To40, males40To50, males50To60,
    ↪ males60To70, males70To80, malesAbove80)
femaleAges = (femalesUnder20, females20To30, males30To40, females40To50,
    ↪ females50To60, females60To70, females70To80, femalesAbove80)
ind = np.arange(N)
width = 0.35

figure, axes = plt.subplots()
figure.set_size_inches(15, 10, forward=True)
plt.bar(ind, maleAges, width, label='Male', color='cornflowerblue')
plt.bar(ind + width, femaleAges, width, label='Female', color='hotpink')


plt.xlabel('Age Groups')
plt.ylabel('Number of Customers')
plt.title('Average Ages of Customers')

plt.xticks(ind + width / 2, ('<20', '20-29',
    ↪ '30-39', '40-49', '50-59', '60-69', '70-79', '80+'))

for p in axes.patches:
    axes.annotate(f'{np.round((p.get_height()/total)*100, decimals=2)}%',
        xy=(p.get_x()+p.get_width()/2., p.get_height()),
        ha='center',
        va='center',
        xytext=(0, 10),
        textcoords='offset points')

plt.legend(loc='best')
plt.show()

```



nhammad

1.7 Visualizing Purchases Based on Industries

```
[9]: print(df.job_industry_category.value_counts(), "\n")

total = df.job_industry_category.size
null_values = sum(pd.isnull(df['job_industry_category']))
total = total-null_values

industries = (df['job_industry_category'].sort_values()).value_counts().keys().
    ↳ tolist()
counts = (df['job_industry_category'].sort_values()).value_counts().tolist()


data = sorted(zip(industries, counts), key=lambda v: v[1], reverse=True)

plt.figure(figsize=(16, 6))
for (i, c) in data:
    bars = plt.bar(i, c, width=0.2, bottom=None, align='center', data=None,
    ↳ color='cornflowerblue')
    plt.annotate(f'{int((c/total)*100)}%\n', xy=(i, c), va='center',
    ↳ ha='center')
plt.xlim(-0.9, len(industries) - 1 + 0.9)
```

```
plt.xlabel('Industry')
plt.ylabel('Number of Customers')
plt.title('Industry Distribution of Customers')
plt.show()
```

Manufacturing	695
Financial Services	682
Health	543
Retail	305
Property	231
IT	126
Entertainment	122
Agriculture	99
Telecommunications	61

Name: job_industry_category, dtype: int64



1.8 Visualizing Purchases Based on Wealth Segments

```
[10]: print(df.wealth_segment.value_counts(), "\n")

total = df.wealth_segment.size
null_values = sum(pd.isnull(df['wealth_segment']))
total = total-null_values

segments = (df['wealth_segment'].sort_values()).value_counts().keys().tolist()
counts = (df['wealth_segment'].sort_values()).value_counts().tolist()

data = sorted(zip(segments, counts), key=lambda v: v[1], reverse=True)
```

```


plt.figure(figsize=(16, 6))
for (i, c) in data:
    bars = plt.bar(i, c, width=0.2, bottom=None, align='center', data=None,
        color='cornflowerblue')
    plt.annotate(f'{int((c/total)*100)}%\n', xy=(i, c), va='center',
        ha='center')
plt.xlim(-0.9, len(segments) - 1 + 0.9)
plt.xlabel('Wealth Segment')
plt.ylabel('Number of Customers')
plt.title('Wealth Segment Distribution of Customers')
plt.show()

```

```

Mass Customer      1695
High Net Worth     871
Affluent Customer  849
Name: wealth_segment, dtype: int64

```



1.9 Visualizing Purchases Based on Whether the Customers Own A Car or Not

```

[11]: print(df.owns_car.value_counts(), "\n")


labels = (df['owns_car'].sort_values()).value_counts().keys().tolist()
counts = (df['owns_car'].sort_values()).value_counts().tolist()

figureObject, axesObject = plt.subplots()
axesObject.pie(counts, labels=labels, autopct='%1.0f',
    colors=['tomato', 'springgreen'])
axesObject.axis('equal')

```

```
plt.title("Do Customers Own Cars or Not? (%)", bbox={'facecolor':'0.8', 'pad':
↪5})
plt.show()
```

```
Yes      1734
No       1681
Name: owns_car, dtype: int64
```



2 Customer Addresses

Now it's time to involve the second dataset.

2.1 Removing Contradiction in States

Upon seeing the head of our dataset, we observe something strange. Some state names are written full while others are abbreviations. To investigate further, we can check the value counts. After checking the value counts, we realize that there's a contradiction in state names. For instance, New South Wales and NSW is the same but are written separately. Same is the case with Victoria and VIC. Hence, we make the necessary replacements.

[12]:

```
customerAddressDF = pd.read_excel("KPMG_VI_New_raw_data_update_final.xlsx",
    ↳sheet_name=4, skiprows=1)
print(customerAddressDF.state.value_counts(), "\n")
customerAddressDF["state"].replace({"New South Wales": "NSW", "Victoria":
    ↳"VIC"}, inplace=True)
```


```
NSW                2054
VIC                 939
QLD                 838
New South Wales     86
Victoria            82
Name: state, dtype: int64
```

2.2 Visualizing Purchases Based on States

```
[13]: states = (customerAddressDF['state'].sort_values()).value_counts().keys().
    ↳tolist()
counts = (customerAddressDF['state'].sort_values()).value_counts().tolist()

data = sorted(zip(states, counts), key=lambda v: v[1], reverse=True)

plt.figure(figsize=(16, 6))
for (i, c) in data:
    bars = plt.bar(i, c, width=0.2, bottom=None, align='center', data=None,
    ↳color='cornflowerblue')
    plt.annotate(f'{int((c/total)*100)}%\n', xy=(i, c), va='center',
    ↳ha='center')
plt.xlim(-0.9, len(states) - 1 + 0.9)
plt.xlabel('State')
plt.ylabel('Number of Customers')
plt.title('Location of Customers')
plt.show()
```

2.3 Visualizing Purchases Based on States & Gender

```
[14]: customerAddressDF = pd.merge(customerAddressDF, df, on='customer_id',  
    ↪how='inner')  
  
maleNSW = 0  
maleVIC = 0  
maleQLD = 0  
  
femaleNSW = 0  
femaleVIC = 0  
femaleQLD = 0  
  
NSWData = customerAddressDF[customerAddressDF['state']=="NSW"]  
VICData = customerAddressDF[customerAddressDF['state']=="VIC"]  
QLDData = customerAddressDF[customerAddressDF['state']=="QLD"]  
  
for value in NSWData['gender']:  
    if(value=="F"):  
        femaleNSW+=1  
    else:  
        maleNSW += 1  
  
for value in VICData['gender']:  
    if(value=="F"):  
        femaleVIC+=1  
    else:  
        maleVIC += 1
```

```

for value in QLDData['gender']:
    if(value=="F"):
        femaleQLD+=1
    else:
        maleQLD += 1

N = 3
labels = ['NSW','VIC', 'QLD']
male = (maleNSW, maleVIC, maleQLD)
female = (femaleNSW, femaleVIC, femaleQLD)
ind = np.arange(N)
width = 0.30

figure, axes = plt.subplots()
figure.set_size_inches(15, 10, forward=True)
plt.bar(ind, male , width, label='Male', color='cornflowerblue')
plt.bar(ind + width, female, width, label='Female', color='hotpink')


plt.xlabel('States')
plt.ylabel('Number of Customers')
plt.title('State & Gender Distribution of Customers')

plt.xticks(ind + width / 2, ('NSW','VIC', 'QLD'))

for p in axes.patches:
    axes.annotate(f'{np.round((p.get_height()/total)*100, decimals=1)}%',
                  xy=(p.get_x()+p.get_width()/2., p.get_height()),
                  ha='center',
                  va='center',
                  xytext=(0, 10),
                  textcoords='offset points')

plt.legend(loc='best')
plt.show()

```



nhammad